



FIT3143 Lab Week 10

Lecturers: ABM Russel (MU Australia) and Vishnu Monn (MU Malaysia)

MPI VIRTUAL TOPOLOGY + COMM SPLIT & PIPELINE COMPUTATION

OBJECTIVES

- The purpose of this lab is to first learn how to combine MPI Virtual Topologies with MPI Comm Split. Then, this lab explores integrating a thread into an MPI process.
- This lab also explores pipeline computation using MPI.

INSTRUCTIONS

- Download and set up the Linux VM [Refer to Lab Week 1]
- Setup eFolio (including Git) and share with tutor and partner [Refer to Lab Week 1]

TASK

DESCRIPTION:

- Combining MPI virtual topology functions and MPI Comm Split.
- Integrating a thread into an MPI process function.
- Design and implement an MPI pipeline computation program.

WHAT TO SUBMIT:

1. Algorithm or code description, analysis of results, screenshot of the running programs and git repository URL in the eFolio.
2. Code in the Git.

EVALUATION CRITERIA

- This Lab-work is part of grading
- Code compile without errors (2), well commented (2), lab-work questions fully answered (4), analysis/report is well formatted (2) = 10 marks

LAB ACTIVITIES (10 MARKS)

Task1 – Placing the Slaves into a virtual topology – Worked example

This task continues from Lab Week 09's Task 4. Here, the slaves are first placed in a 2D virtual topology (using MPI Cartesian functions). Each slave then sends a series of messages to the Master, which prints the message.

Sample solution:

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <mpi.h>
#include <time.h>
#include <unistd.h>

#define MSG_EXIT 1
#define MSG_PRINT_ORDERED 2
#define MSG_PRINT_UNORDERED 3

int master_io(MPI_Comm world_comm, MPI_Comm comm);
int slave_io(MPI_Comm world_comm, MPI_Comm comm);

int main(int argc, char **argv)
{
    int rank, size;
    MPI_Comm new_comm;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    MPI_Comm_split( MPI_COMM_WORLD, rank == size-1, 0, &new_comm);
    // color will either be 0 or 1
    if (rank == size-1)
        master_io( MPI_COMM_WORLD, new_comm );
    else
        slave_io( MPI_COMM_WORLD, new_comm );
    MPI_Finalize();
    return 0;
}

/* This is the master */
int master_io(MPI_Comm world_comm, MPI_Comm comm)
{
    int i, size, nslaves, firstmsg;
    char buf[256], buf2[256];
```

```

MPI_Status status;
MPI_Comm_size(world_comm, &size );
nslaves = size - 1;

while (nslaves > 0) {
    MPI_Recv(buf, 256, MPI_CHAR, MPI_ANY_SOURCE,
MPI_ANY_TAG, world_comm, &status );
    switch (status.MPI_TAG) {
        case MSG_EXIT: nslaves--; break;
        case MSG_PRINT_UNORDERED:
            fputs( buf, stdout );
            break;
        case MSG_PRINT_ORDERED:
            firstmsg = status.MPI_SOURCE;
            for (i=0; i<nslaves; i++) {
                if (i == firstmsg)
                    fputs( buf, stdout );
                else {
                    MPI_Recv( buf2, 256, MPI_CHAR, i,
MSG_PRINT_ORDERED, world_comm, &status );
                    fputs( buf2, stdout );
                }
            }
            break;
    }
}
return 0;
}

/* This is the slave */
int slave_io(MPI_Comm world_comm, MPI_Comm comm)
{
    int ndims=2, size, my_rank, reorder, my_cart_rank, ierr,
worldSize;
    MPI_Comm comm2D;
    int dims[ndims], coord[ndims];
    int wrap_around[ndims];
    char buf[256];

    MPI_Comm_size(world_comm, &worldSize); // size of the world
communicator
    MPI_Comm_size(comm, &size); // size of the slave communicator
    MPI_Comm_rank(comm, &my_rank); // rank of the slave
communicator
    dims[0]=dims[1]=0;

    MPI_Dims_create(size, ndims, dims);
    if(my_rank==0)
        printf("Slave Rank: %d. Comm Size: %d: Grid Dimension =
[%d x %d] \n",my_rank,size,dims[0],dims[1]);

    /* create cartesian mapping */

```

```

wrap_around[0] = 0;
wrap_around[1] = 0; /* periodic shift is .false. */
reorder = 0;
ierr = 0;
ierr = MPI_Cart_create(comm, ndims, dims, wrap_around,
reorder, &comm2D);
if(ierr != 0) printf("ERROR[%d] creating CART\n",ierr);

/* find my coordinates in the cartesian communicator group */
MPI_Cart_coords(comm2D, my_rank, ndims, coord); //
coordinated is returned into the coord array
/* use my cartesian coordinates to find my rank in cartesian
group*/
MPI_Cart_rank(comm2D, coord, &my_cart_rank);

/*
printf("Global rank (within slave comm): %d. Cart rank: %d.
Coord: (%d, %d).\n", my_rank, my_cart_rank, coord[0], coord[1]);
fflush(stdout);
*/

sprintf( buf, "Hello from slave %d at Coordinate: (%d,
%d)\n", my_rank, coord[0], coord[1]);
MPI_Send( buf, strlen(buf) + 1, MPI_CHAR, worldSize-1,
MSG_PRINT_ORDERED, world_comm );

sprintf( buf, "Goodbye from slave %d at Coordinate: (%d,
%d)\n", my_rank, coord[0], coord[1]);
MPI_Send(buf, strlen(buf) + 1, MPI_CHAR, worldSize-1,
MSG_PRINT_ORDERED, world_comm);

sprintf(buf, "Slave %d at Coordinate: (%d, %d) is exiting\n",
my_rank, coord[0], coord[1]);
MPI_Send(buf, strlen(buf) + 1, MPI_CHAR, worldSize-1,
MSG_PRINT_ORDERED, world_comm);
MPI_Send(buf, 0, MPI_CHAR, worldSize-1, MSG_EXIT,
world_comm);

MPI_Comm_free( &comm2D );
return 0;
}

```

Task2 – Adding a thread as an asynchronous activity at the Master – Worked Example

This task is based on modifying the sample solution code from Task 1 above. Move the **bolded region of the while loop code** in the sample solution of Task 1 into a thread function. This means that the `master_io()` function does the following:

- a) Creates a thread
- b) Waits for the thread to complete
- c) Exits

The thread function implements the bolded while loop code (i.e., waiting for messages from the slaves and printing these messages). Make sure to pass the necessary values from the `master_io()` function to the thread.

You can use the POSIX thread library to create the thread here. You could also consider using OpenMP as an asynchronous thread.

Note: You may feel that this task is unnecessary as the Task 1 program code works just fine. However, the aim of Task 2 here is to provide you some basic exposure on using thread with MPI.

Sample solution:

```
// Sample solution focuses on the thread function and master_io function.
```

```
void* ProcessFunc(void *pArg) // Common function prototype
{
    int i = 0, size, nslaves, firstmsg;
    char buf[256], buf2[256];
    MPI_Status status;
    MPI_Comm_size(MPI_COMM_WORLD, &size );

    int* p = (int*)pArg;
    nslaves = *p;

    while (nslaves > 0) {
        MPI_Recv(buf, 256, MPI_CHAR, MPI_ANY_SOURCE,
MPI_ANY_TAG, MPI_COMM_WORLD, &status );
        switch (status.MPI_TAG) {
            case MSG_EXIT: nslaves--; break;
            case MSG_PRINT_UNORDERED:
                printf("Thread prints: %s", buf);
                fflush(stdout);
                break;
            case MSG_PRINT_ORDERED:
                firstmsg = status.MPI_SOURCE;
                for (i=0; i<size-1; i++) {
                    if (i == firstmsg){
                        printf("Thread prints: %s", buf);
                        fflush(stdout);
                    }else {
```

```

        MPI_Recv( buf2, 256, MPI_CHAR, i,
MSG_PRINT_ORDERED, MPI_COMM_WORLD, &status );
        printf("Thread prints: %s", buf2);
        fflush(stdout);
    }
}
break;
}
}

return 0;
}

int master_io(MPI_Comm world_comm, MPI_Comm comm)
{
    int size, nslaves;
    MPI_Comm_size(world_comm, &size );
    nslaves = size - 1;

    pthread_t tid;
    pthread_create(&tid, 0, ProcessFunc, &nslaves); // Create the
thread
    pthread_join(tid, NULL); // Wait for the thread to complete.

    return 0;
}

```

Task3 – Simple pipeline computing using MPI – Worked example

A series of raw experimental results have been stored into the **ExpResults.txt** file, as shown in Figure T3. The first element in **ExpResults.txt** represents the number of experiment results in this file. Subsequent elements in this file represent the respective experiment results.




Figure T3: Content of *ExpResult.txt*

Each piece of result from this file, x_0 , has been subjected to further calculations as per the following mathematical expressions:

$$x_1 = x_0 - 4x_0 + 7$$

$$x_2 = x_1^3 + \sin\left(\frac{x_1}{8}\right)$$

$$x_3 = 2x_2^4 + \cos(4x_2) + 3\pi$$

$$x_4 = 3x_3^2 - 2x_3 + \frac{\tan(x_3)}{3}$$

Using the C programming language with a parallel programming implementation using the Message Passing Interface (MPI):

Write a program to apply the content of **ExpResults.txt** into the series of equations as aforementioned such that the result for each data content is represented by x_4 .

Only the root rank is permitted to access the data content of **ExpResults.txt** and read each datum one by one. Each datum read from this file is stored into variable x_0 . The root node then calculates x_1 based on the value of x_0 . The result of x_1 is then transmitted to the subsequent node, which calculates x_2 . This process continues in a **parallel pipeline structure** until the last node calculates x_4 .

In addition, only the root node will print out the results of x_4 for each datum read from file **ExpResults.txt**. To achieve this design, the last node will store a list of calculated x_4 results before sending this list back to the root node for result printout. Use dynamic memory allocation to create a list based on the number of elements in the file.

Note: The total number of nodes in the pipeline architecture should be four, with x_1, x_2, x_3 & x_4 representing the first, second, third and fourth nodes respectively. Please create your own ExpResults.txt file based on the format as seen in Figure T3. You need not create a large number of entries for your version of ExpResults.txt. A smaller number of entries would suffice to test the program.

Sample solution:

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <memory.h>
#include "mpi.h"

#define SENTINEL 0.0

int main(int argc, char *argv[])
{
    FILE *pInfile;
    double x0, x1, x2, x3, x4;
    double *pX4Buff = NULL;
    float x;
    int fileElementCount = 0;
    int counter = 0;

    int my_rank;
```

```

int p;
MPI_Status status;

MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
MPI_Comm_size(MPI_COMM_WORLD, &p);

switch (my_rank)
{
    case 0:
    {
        pInfile = fopen("ExpResults.txt","r");
        fscanf(pInfile, "%d", &fileElementCount);

        pX4Buff = (double*)malloc(fileElementCount *
sizeof(double));
        memset(pX4Buff, 0, fileElementCount *
sizeof(double));

        // Send the counter to the last process
        MPI_Send(&fileElementCount, 1, MPI_INT, (p - 1), 0,
MPI_COMM_WORLD);

        // Read each element from the file
        while(counter < fileElementCount)
        {
            fscanf(pInfile, "%f", &x);
            x0 = x;
            x1 = x0 - (4 * x0) + 7;
            MPI_Send(&x1, 1, MPI_DOUBLE, 1, 0,
MPI_COMM_WORLD);
            counter++;
        }

        // File end, send a SENTINEL value to complete
        calculation
        fclose(pInfile);
        pInfile = NULL;

        x1 = SENTINEL;
        MPI_Send(&x1, 1, MPI_DOUBLE, 1, 0, MPI_COMM_WORLD);

        // Wait for buffer from last node
        MPI_Recv((void*)pX4Buff, counter, MPI_DOUBLE, (p -
1), 0, MPI_COMM_WORLD, &status);

        // Print results
        for(int i = 0; i < counter; i++)
        {
            printf("Result[%d]: %g\n", i, pX4Buff[i]);
        }
        free(pX4Buff);
    }
}

```



```

        pX4Buff = NULL;
        break;
    }
    case 1:
    {
        do
        {
            MPI_Recv(&x1, 1, MPI_DOUBLE, 0, 0,
MPI_COMM_WORLD, &status);
            if(x1 != SENTINEL)
            {
                x2 = pow(x1, 3) + sin(x1/8);
            }
            else
            {
                x2 = SENTINEL;
            }
            MPI_Send(&x2, 1, MPI_DOUBLE, 2, 0,
MPI_COMM_WORLD);
        } while (x1 != SENTINEL);
        break;
    }
    case 2:
    {
        do
        {
            MPI_Recv(&x2, 1, MPI_DOUBLE, 1, 0,
MPI_COMM_WORLD, &status);
            if(x2 != SENTINEL)
            {
                x3 = (2 * pow(x2, 4)) + cos(4 * x2) + (3
* M_PI);
            }
            else
            {
                x3 = SENTINEL;
            }
            MPI_Send(&x3, 1, MPI_DOUBLE, 3, 0,
MPI_COMM_WORLD);
        } while (x2 != SENTINEL);
        break;
    }

    case 3:
    {
        // Get the file element count first
        MPI_Recv(&fileElementCount, 1, MPI_DOUBLE, 0, 0,
MPI_COMM_WORLD, &status);
        pX4Buff = (double*)malloc(fileElementCount *
sizeof(double));
        memset(pX4Buff, 0, fileElementCount *
sizeof(double));
    }

```

```
// Now, receive the pipelined data
counter = 0;
do
{
    MPI_Recv(&x3, 1, MPI_DOUBLE, 2, 0,
MPI_COMM_WORLD, &status);
    if(x3 != SENTINEL)
    {
        x4 = (3 * pow(x3, 2)) - (2 * x3) +
(tan(x3) / 3);

        // Save the result into buffer &
increment the buffer counter
        pX4Buff[counter] = x4;
        counter++;
    }
} while (x3 != SENTINEL);

// End of file reached, send the buffer back to the
root
if(counter > 0)
{
    MPI_Send(pX4Buff, counter, MPI_DOUBLE, 0, 0,
MPI_COMM_WORLD);
}
// Clean up
free(pX4Buff);
pX4Buff = NULL;
break;
}
default:
{
    printf("Process %d unused.\n",my_rank);
    break;
}
}

MPI_Finalize();

return 0;
}
```

Task 4 – Pipeline computing exercise using MPI

In mathematics, a quadratic equation represents a univariate polynomial equation of the second degree. A general quadratic equation can be described as:

$$ax^2 + bx + c = 0 \quad (1.1)$$

where x represents the unknown variable and a , b and c are the quadratic coefficients ($a \neq 0$). A quadratic equation with real and complex coefficients has two solutions, called roots (x_1 and x_2).

The discriminant, d , is computed as: $d = b^2 - 4ac$. If d is positive, the quadratic equation has two distinct real roots (i.e., $x_1 \neq x_2$) such that:

$$x_1 = \frac{-b + \sqrt{d}}{2a}, x_2 = \frac{-b - \sqrt{d}}{2a} \quad (1.2)$$

If d is zero, the quadratic equation has only one real root (i.e., $x_1 = x_2$) such that:

$$x_1 = x_2 = \frac{-b}{2a} \quad (1.3)$$

If d is negative, the quadratic equation has two distinct complex roots (i.e., $x_1 \neq x_2$) such that:

$$x_1 = \frac{-b}{2a} + \frac{i\sqrt{|d|}}{2a}, x_2 = \frac{-b}{2a} - \frac{i\sqrt{|d|}}{2a} \quad (1.4)$$

Figure T4-1 illustrates the content of a text file, *quad.txt*, which contains a set of quadratic coefficients. The first row element in *quad.txt* represents the number of coefficients (a , b and c) rows in this file. The second row displays the legend text for these coefficients and the third row onwards contains the coefficient content.




Figure T4-1: Content of *quad.txt*

Using the C programming language with a parallel programming implementation using the Message Passing Interface (MPI):

- a) Write a program to compute the quadratic roots of each row coefficient in *quad.txt* file.
- Only the root rank (or first node) is permitted to access the *quad.txt* file.
 - The root rank reads coefficients of each row from *quad.txt* file and computes the discriminant, d , which is then transmitted to the subsequent node along with the a and b coefficients.
- b) Continuing from part (a), the second node receives the computed d and a and b coefficients (per row) from the first node and computes the root values.
- This node computes x_1 and x_2 based on the computed value of d .
 - Note: If $d < 0$, the roots are calculated as x_{1_real} , x_{1_img} and x_{2_real} , x_{2_img} .
 - The computed root values are then transmitted to the third node.
- c) Continuing from part (b), the third node receives the computed root values from the second node and writes these root values into a new text file, *roots.txt*. Figure T4-2 illustrates a sample content of the computed root value in the *roots.txt* file.

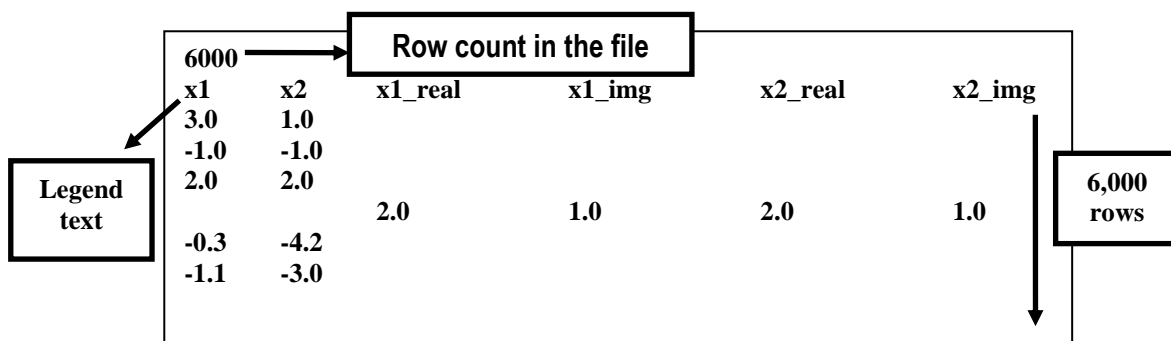


Figure T4-2: Content of *roots.txt*

Complete parts (a), (b) and (c). Use a **parallel pipeline structure** in reading the coefficients from the file, computing the roots and writing the computed roots into file.

Note: The total number of nodes in this pipeline architecture should be three (The first node reads a row coefficient and computes d , the second node computes the roots and the third node writes the computed roots into a new file). Create your own version of *quad.txt* as seen in Figure T4-1. You need not create a large number of entries for your version of *quad.txt*. A smaller number of entries would suffice to test the program.

You may refer to the following C code to get you started.

```
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <mpi.h>

int main()
{
    FILE *pInfile;
    float a_coeff, b_coeff, c_coeff, x1, x2, disc;
    float x1r, x1i, x2r, x2i;
    int fileElementCount = 0, constNeg = -1;;

    int my_rank;
    int p;
    MPI_Status status;

    MPI_Init(&argc, &argv);
```

```
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
MPI_Comm_size(MPI_COMM_WORLD, &p);

// WRITE PART(a) HERE

switch (my_rank){
    case 0:{
        // CONTINUE WITH PART (a) HERE
        break;
    }
    case 1:{
        // WRITE PART (b) HERE
        break;
    }
    case 2:{
        // WRITE PART (c) HERE
        break;
    }
}
MPI_Finalize();
return 0;
}
```

OPTIONAL BONUS ACTIVITIES (2 MARKS)

Develop a pipeline solution to compute *sin* according to

$$\sin \theta = \theta - \frac{\theta^3}{3!} + \frac{\theta^5}{5!} - \frac{\theta^7}{7!} + \frac{\theta^9}{9!} - \dots$$

A series of values represent the input, $\theta_0, \theta_1, \theta_2, \theta_3, \dots$